

A Framework for Emergent Time from Horn-Filling Dynamics

Russi Chatterjee

August 2, 2026

Abstract

We propose a framework in which spacetime is a growing simplicial set and time is the process of filling its inner horns. Two local, non-deleting moves drive the growth of spacetime: a "sprout" attaches a fresh directed edge ending at a new vertex, and a "fill" completes an unfilled inner horn with a composite edge and its witnessing cell. Each move is a weak homotopy equivalence of the underlying space, so the homotopy type is conserved at every step. Since the geometric realization of the simplicial set forgets edge orientation, this conserved type cannot see any arrow of time; all physical content lives in the directed structure that realization forgets. The completed limit of the filling process is a quasicategory, in which every composition question has an answer while edges remain directed. Each move is an event that consumes cells created by earlier events, the dependency of moves forms a partial order, and a reference frame is a choice of linearization of this order. The entropy of an event is the logarithm of the number of linear extensions of its causal past. This quantity measures how much of the past could have happened in a different order, and grows monotonically, which we prove; a toy implementation provides first numerical support. The arrow of time appears as a broken symmetry. The universe carries an arrow exactly when it is not isomorphic to its own edge-reversal, and pairing the universe with its mirror image restores the symmetry globally, in the spirit of CPT-symmetric cosmologies. We formulate a maximum-entropy inference problem whose solution, conjecturally of Gibbs form, would determine the rule from locality and a fixed update budget alone. We close with the open problems, foremost the recovery of the Lorentz group and of an area law for horizon entropy.

1 Introduction

Discrete approaches to quantum gravity, from Regge calculus [14] to causal dynamical triangulations [3], replace smooth spacetime with combinatorial data: simplices, gluings, and the lengths assigned to them. In Regge calculus the combinatorial scaffold is fixed with only edge lengths allowed to vary. The vertices, simplices, and gluings of the triangulation never change, so no new combinatorial structure is ever introduced. That is enough for classical gravity, where one solves the field equations on a triangulation that is given in advance. The discomfort begins one step later, when Feynman's path integral is applied to the evolution of the universe and the sum runs over all histories. Summing over Regge geometries means integrating over edge lengths, a continuous space in which many assignments describe the same geometry and no canonical measure presents itself. On the other hand, dynamical triangulations [8, 1, 2] escape this by inverting Regge's choice. The lengths are frozen and the combinatorics is set free, so a geometry becomes a gluing pattern of simplices whose shapes are fixed once and for all, and summing over geometries becomes counting gluings.

In both cases the geometry is carried by the gluing and not by any ambient space. Each simplex is flat inside, so curvature can only sit where simplices meet, and it is measured by how far the angles around such a meeting place fall short of a full turn. Figure ?? shows the three cases for triangles glued around a vertex.

Because the combinatorics is now free, the number of simplices differs from history to history, and within a single history the number of simplices in a slice changes from one slice to the next.




Figure 1: Curvature in Regge calculus lives at the vertices, not on the interiors of the simplices. Gluing k flat unit equilateral triangles around a single point makes the angles there sum to $k \cdot 60^\circ$, and the deficit angle $\delta = 2\pi - \sum_i \theta_i$ measures the failure to close up in the plane. Top row shows the resulting surface in $\mathbb{R}^3$, bottom row shows the same triangles unfolded flat, where δ is directly visible. For $k = 5$ the angles fall short, the gap closes only by lifting the vertex out of the plane, and the cone carries positive curvature concentrated at a point. For $k = 6$ the angles close exactly and the neighbourhood is isometric to a patch of the Euclidean plane. For $k = 7$ the seventh triangle (dashed) overlaps the first, the excess forces the surface to buckle into a saddle, and the curvature is negative. Every triangle is flat throughout, so all the geometry is carried by the combinatorics of how they are glued.

New geometric measure appears from nowhere. Classical general relativity is no better off on this point, since it has no conservation law for spatial volume either, and in a closed universe the Hamiltonian constraint balances the total at zero [4, 9] however much volume appears. The complaint is therefore not that volume grows, but that the growth answers to nothing the theory conserves. A cosmological constant term only assigns a cost to four-volume within the ensemble, and where the total is held fixed instead, as in simulations and in unimodular gravity [16, 11], the constraint is imposed by hand rather than descending from a symmetry.

This paper takes the discomfort seriously and asks what a theory looks like if one demands, from the start, that

1. Nothing is ever destroyed. Evolution proceeds only by forking and completion moves, never by deletion, so that topological invariants are conserved by construction.
2. Time is not a dimension of the fundamental structure but an emergent bookkeeping of an irreversible forking-filling process.
3. Entropy is attached to frame-invariant causal data, not to spatial slices, so that the theory never needs to say what all of space looks like at a single moment.
4. The rule that chooses which move happens next is derived from an inference principle rather than postulated.

In our framework the universe begins as a single vertex, which we call the big bang, and grows by two kinds of local moves:

1. A **sprout** attaches a fresh directed edge ending at a new vertex, adding new space.

2. A **fill** completes an unfilled inner horn, attaching a composite edge together with the higher cell that witnesses the composition.

Two kinds of direction appear here. Every edge of the simplicial set carries an orientation, and that orientation is what decides which horns count as inner and what composition means. Separately, every move is an event, and the events are ordered by which of them created the cells the others consumed. We call the first the **orientation** and the second the **causal order**. No move ever deletes anything, so growth always runs forward in the causal order, while orientation is a property of the finished cells and can be reversed on paper.

At the big bang two edges sprout from the single vertex with opposite orientation, and this is what produces the two branches. The mirror branch is not an independent run. By definition it is the edge-reversal of the branch, carrying one mirrored event for every event, which is what makes the pair a mirror image of itself. The mirroring is causally inert, since no move in one branch ever consumes a cell created in the other, so the causal past of any event lies wholly inside its own branch. Each branch grows forward in its own causal order while pointing the opposite way in orientation, so neither branch runs backward in time in any absolute sense. Each is backward only as seen from the other, and there is no global causal order in which either of them is reversed. This is a combinatorial counterpart of CPT-symmetric cosmologies, where a universe and an antiuniverse emerge together from a single point [6, 5].




Figure 2: The total object. A branch on the right and its mirror image on the left, glued at the single vertex where both were born. Corresponding edges carry opposite orientation, so the pair is isomorphic to its own edge-reversal. The shaded region on each side is a filled inner horn, drawn with its composite edge in bold, and the fill is mirrored along with everything else. The grey arrows are the causal order, which runs away from the shared vertex on both sides, so each branch is backward only as seen from the other, and never in the order in which its own cells were built.

At any moment the state of the universe is a finite simplicial set with directed edges, together with the history of how it was built. The limit of the filling process in one branch is a quasicategory, in which every composition has been formed while the orientations remain. The classical, zero-information limit is the nerve of an ordinary category, the simplicial set whose n -cells are the chains of n composable morphisms [15, 12], and in it every filler is unique. The physical content of a completed universe is precisely how far it is from being a nerve.

The pair of branches is meant to carry no information at all. Two contractible halves glued at a single vertex are again contractible, so every homotopy invariant of the pair is that of a point, and the pair is isomorphic to its own edge-reversal, so it carries no arrow either. Whatever is physical therefore sits inside a single branch, in the orientations and in the history of how its cells were built.

2 Preliminaries and related work

3 The framework

3.1 The object

At any moment, the universe is a finite simplicial set with directed edges. It begins as a single vertex Δ^0 , the big bang, and it grows from there.

The object is pre-geometric. It is not spacetime, and it is not a representation of a spacetime lying elsewhere. It carries no metric, no dimension, and no manifold structure. Everything geometric has to be recovered from it. Statements about space are made in the framework's own combinatorial terms.

3.1.1 Growing the Simplicial Set

Definition 3.1. The simplicial set grows by two local, non-deleting **moves**, which are pushouts along the following two monomorphisms of finite simplicial sets:

1. **Sprout:** $\Delta^0 \hookrightarrow \Delta^1$. A sprout can happen at any vertex v of the current state. Performing it attaches a fresh directed edge from v to a new vertex.
2. **Fill:** $\Lambda_k^n \hookrightarrow \Delta^n$ for $0 < k < n$. A fill can happen at any unfilled inner horn. Performing it attaches the missing face and interior. Each horn may be filled at most once.




Figure 3: The two moves and what each of them adds. In (a) a vertex sprouts a fresh directed edge to a new vertex. After a second sprout, (b) shows the composable pair $v_0 \rightarrow v_1 \rightarrow v_2$ whose composite is absent, an unfilled inner horn Λ_1^2 (amber, missing edge dashed). In (c) the fill supplies the composite $v_0 \rightarrow v_2$ (green) together with the 2-cell witnessing it (shaded). Panel (d) shows the state after a third sprout and both first-tier fills. Two distinct horns now call for a composite from v_0 to v_3 , one built from $v_0 \rightarrow v_1$ with $v_1 \rightarrow v_3$ and one from $v_0 \rightarrow v_2$ with $v_2 \rightarrow v_3$. Each fill mints its own composite, so a and b are parallel edges rather than a single shared one. Only once both are present does the inner horn Λ_1^3 exist, since its three faces already carry all six edges of Δ^3 ; filling it therefore adds a 2-cell and a 3-cell but no edge. Only inner horns are ever filled, so edge direction is never inverted and the limit object is a quasicategory.

Remark 3.2. The initial object is a single vertex Δ^0 (the big-bang event). At any moment the current state is a finite simplicial set with directed edges, and the next move is chosen from the set of available moves, which is the disjoint union of all sprout sites and all unfilled inner horns.

Remark 3.3. Every vertex is created by exactly one sprout, which supplies its first incoming edge. Every other incoming edge is a composite supplied by a fill. So sprout edges and composite edges never mix. A sprout creates a fresh vertex and attaches a directed edge to it. The composite joins two vertices that already exist.

Remark 3.4. Sprouting is a Whitehead elementary expansion and inner-horn filling is an anodyne extension. Both operations preserve homotopy equivalences, so the homotopy type of the simplicial set is conserved by construction. In our case, the homotopy type remains that of a point.

Remark 3.5. Only two moves ever add an edge. A sprout adds one, and a fill in dimension two adds one. Fills in dimension three and above add no edges at all, only cells that say existing paths agree. So every chain is built out of edges, and nothing above dimension one can be part of a chain. Whatever geometric and metric structure this model has is carried by the vertices, the edges, and the causal order between events. Treat the edges as a multigraph. If two fills each produce a composite, keep both as separate edges. The number of parallel edges is a record of how many fills happened there. The order says which events can affect which. The edge counts say how much happened. Both are needed, and neither replaces the other.

3.2 Events, runs, and the causal order

Growth happens one move at a time. A **run** is a finite partially ordered set of move applications starting from the single vertex. Since no move deletes anything, an event can only consume cells that already exist. Event b therefore depends on event a when the match of b uses a cell created by a , and the reflexive-transitive closure of dependency is the causal order. A reference frame is a linearization of this order.

Definition 3.6. A **run** is a finite partially ordered set of move applications starting from the single vertex Δ^0 . Each application is an event, which may be a sprout or a fill.

At any moment, a run is a partially ordered list of events representing a simplicial set. Formally, write a run as a pair $(E, \leq)$, where $E = \{e_1, \dots, e_n\}$ is the set of events and $\leq$ is the causal order on them.

Example 3.7. Start from the single vertex v_0 . Event s_1 sprouts an edge $v_0 \rightarrow v_1$, event s_2 sprouts an edge $v_1 \rightarrow v_2$, and event s_3 sprouts a second edge $v_0 \rightarrow v_3$ from the starting vertex. Event f fills the horn $v_0 \rightarrow v_1 \rightarrow v_2$, supplying the composite $v_0 \rightarrow v_2$ and its witnessing 2-cell. Event s_4 sprouts an edge $v_2 \rightarrow v_4$, and event g fills the horn $v_0 \rightarrow v_2 \rightarrow v_4$, consuming the composite from f . The causal order is generated by $s_1 < s_2 < f < g$ and $s_2 < s_4 < g$, while s_3 shares only the initial vertex with the others and stays incomparable to all of them. The causal past of g is $J^-(g) = \{s_1, s_2, f, s_4, g\}$, in which f and s_4 are incomparable, so $\text{Ext}(J^-(g)) = 2$ and $S(g) = \log 2$. Note that the state at g does not contain v_3 , since s_3 is not in the causal past of g .




Figure 4: Example 3.7. Left: the simplicial set after the six events, with the composite edges supplied by the fills f (green) and g (orange) and their witnessing 2-cells shaded to match. Right: the causal order as a Hasse diagram. The events f and s_4 are incomparable, which is what gives $J^-(g)$ its two linearizations. The event s_3 and the vertex v_3 it created (grey) lie outside the causal past of g , so neither enters the state at g .

Definition 3.8. Let e be an event of a run, with causal past $J^-(e) = \{a : a \leq e\}$. The **state at e** is the colimit of the diagram of rule applications in $J^-(e)$.

The colimit is used here so that ordering of $J^-(e)$ has no consequence on the resulting state. Whether the diagram is well-defined, and whether it agrees with what any sequential build produces, is the content of Proposition 3.9.

Proposition 3.9. *The causal order on the events of a run is a partial order, and the colimit in Definition 3.8 exists and agrees with the state produced by any sequential build along a linearization of the run.*

Proof. Write $(E, \leq)$ for the run and X_D for the colimit of the diagram of rule applications over a downset $D \subseteq E$, which exists whenever the diagram is defined since $\mathbf{sSet}$ is cocomplete. A move application at b comes with a match $m_b : L \rightarrow X_{\{a : a < b\}}$, so every cell that b consumes is created by an event strictly below b . Hence $a \prec b$ implies $a < b$, so dependency is a subrelation of a finite partial order and therefore acyclic. Its reflexive-transitive closure is then a partial order, which is the causal order.

For the second claim, define X_D by recursion on the size of the downset D . If $D = \emptyset$ then $X_D = \Delta^0$. Otherwise pick a maximal element b of D , set $D' = D \setminus \{b\}$, and define X_D as the pushout of the rule $L \hookrightarrow R$ of b along $m_b : L \rightarrow X_{D'}$. The match exists by the admissibility condition above, since every cell in the image of m_b is created by an event of D' . The poset is finite, so the recursion terminates and defines the diagram on all downsets of $J^-(e)$, and $X_{J^-(e)}$ is its colimit.

For the third claim, fix a linearization of $J^-(e)$ and consider the sequential build along it. Take two events a and b that are adjacent in the linearization and incomparable in the causal order, applied to a state X . Since a and b are incomparable, neither match uses a cell created by the other, so both matches factor through X . Pushouts of monomorphisms in $\mathbf{sSet}$ are monomorphisms, so the image of m_b persists in $X \cup_{L_a} R_a$ and b can be applied after a . Both orders compute the colimit of the finite diagram formed by X , the two matches, and the two rules, so the two resulting states are canonically isomorphic. This is the local Church–Rosser property for non-deleting double-pushout rewriting [10]. Any two linearizations of a finite poset are connected by a sequence of adjacent transpositions of incomparable elements, so all sequential builds of $J^-(e)$ agree. Finally, a linearization of $J^-(e)$ is a chain containing every object of the poset diagram, so its inclusion into the poset diagram is a final functor, and the iterated pushout along the chain computes the same colimit. Hence every sequential build produces the state of Definition 3.8. $\square$

3.3 Axioms

With the object, the moves, and the causal order in hand, the assumptions of the framework can be stated in the vocabulary just built.

Axiom 1 (Ontology). The fundamental object is the run, a finite partially ordered set of non-deleting moves starting from a single vertex. The state of the universe at an event is the colimit of that run's causal past (Definition 3.8), so it is determined by the run rather than being independent data.

Axiom 2 (Dynamics). Evolution proceeds by two local, non-deleting moves: *sprout* and *fill* (Definition 3.1). Only inner horns may be filled.

Axiom 3 (Budget). The rate of moves is fixed. The constant c is the exchange rate between translation and internal filling (Section 3.6).

Remark 3.10 (What the budget axiom leaves open). A rate requires a clock, and time in this framework is emergent, a frame being only a linearization of the causal order. Read as a universal number of moves per tick, the axiom would reinstate a preferred linearization. The intended reading is the energetic one of Margolus and Levitin [13], in which the budget bounds distinguishable changes per unit energy rather than per unit of any global time, and it awaits a notion of energy internal to the framework. The axiom is also silent on whether the budget is one number for the whole universe or one number per region. Locality favours the latter, and the two readings give different cosmologies (Section 3.8).

Axiom 4 (Inference). The probability distribution over available moves is the maximum-entropy distribution consistent with Axioms 2–3 and locality (Section 3.7). No further dynamical input is allowed.

3.4 The entropy functional

Definition 3.11 (Causal-past entropy). For an event e of a run,

$$S(e) = \log \text{Ext}(J^-(e)),$$

the logarithm of the number of linear extensions of its causal past.

The coarse-graining implicit in S is canonical to the framework. The state, being a colimit, forgets the order of causally independent events, and S counts exactly the forgotten orderings. A chain has $S = 0$. An antichain of n events has $S = \log n!$. Entropy measures causal parallelism.

Lemma 3.12 (Monotonicity). *If $e \leq e'$ then $S(e) \leq S(e')$.*

[GAP: Proof to be written (owner: ixaxaar; the injection of extension sets should be a half-page). Empirically verified in Section ??.]

Lemma 3.13 (Foliation invariance). *$S(e)$ depends only on the isomorphism class of the poset $J^-(e)$ and in particular is independent of any linearization of the run.*

[GAP: Proof to be written; content is the frame-independence of $J^-(e)$.]

Remark 3.14 (Cost of computation). Counting linear extensions is #P-complete in general [7]; Section ?? uses exact counting on small runs and future work requires a sampling estimator.

3.5 Space, distance, and speed

The axioms so far say nothing about space, and no statement about speed has content until they do. This section fixes the minimum, in order.

Definition 3.15 (Weighted one-skeleton). Every edge created by a sprout has weight 1. Every composite edge created by a fill has weight equal to the total weight of the path it composes. The *distance* between two vertices is the least total weight of an undirected path between them in the one-skeleton.

Remark 3.16 (Fills are metrically inert). A composite edge joins two vertices that were already joined by a path of exactly that weight, so under Definition 3.15 a fill changes no distance and creates no reachability that was not already present. The convention is forced if distance is to be stable. Weighting every composite by 1 instead makes each fill a shortcut, and accumulated shortcuts collapse the diameter of the complex in the manner of a small-world graph, leaving neither a stable distance function nor a bounded signal speed. The price of Definition 3.15 is that fills contribute nothing geometric. What they contribute is coherence data, which is where this framework locates its physical content.

Definition 3.17 (Track; translation; internal evolution). A *track* is a chain in the causal order whose consecutive events are matched at cells within bounded distance of one another. An event on a track is a *translation* if it is a sprout, extending the track's locus by one unit of distance, and *internal* if it is a fill, leaving the locus where it is.

Proposition 3.18 (A maximal signal speed). *Along any track, distance advances by at most one unit per event, so the speed of propagation of causal influence is bounded by one unit of distance per event. We write c for this bound.*

[GAP: Proof to be written. Content: by Remark 3.16 a fill adds no reachability, and a sprout adds one vertex at distance 1 from its match.]

The bound of Proposition 3.18 is not derived from anything deep. It is the statement that the moves are local, and c is a unit rather than a prediction. What is not settled, and is the content of problem O1, is why the resulting light-cone structure should be isotropic and the same for every observer.

3.6 Kinematics: proper time and mass

Definition 3.19 (Proper time). The proper time between two causally related events is the length of the longest chain in the causal order between them.

This is the causal set definition of discrete proper time. It is frame-independent, since the longest chain is a property of the poset alone. A null track, all of whose events are translations, has proper time zero, which recovers the statement that a photon carries no internal clock, and conversely a system with internal evolution accumulates proper time and is massive. Mass is the rest-frame rate of internal fills, in the spirit of the Compton clock $mc^2/\hbar$.

Remark 3.20 (Why longest chains, and what dilation still costs). Counting a track's own events gives the wrong dilation law. If a track has N events of which k are internal, the budget splits linearly and the internal rate is $1 - v/c$, where relativity requires $\sqrt{1 - v^2/c^2}$. The two disagree by a factor of three at $v = 0.8c$. The quadratic form comes from the Euclidean geometry of the light-clock argument rather than from any bookkeeping over moves, which is why Appendix ?? states that composition as an assumption. Definition 3.19 is chosen because it is the definition for which the quadratic form is a theorem. For a sprinkling into flat spacetime the longest chain between two events is proportional to $\sqrt{t^2 - x^2}$, and the straight route maximises it. What that theorem assumes is that the causal order resembles a sprinkling into a Lorentzian manifold, so dilation in this framework rests on problem O1 and is not established independently of it.

One further constraint the framework must respect. Uniform motion has to be locally undetectable, so all internal processes of a moving subsystem must slow together and rate differences may only appear on reunion. Any locally measurable change of internal rate under uniform motion would falsify the framework at once.

3.7 The maximum-entropy problem

Definition 3.21 (Move ensemble). At a given state, let $M = M_{\text{sprout}} \sqcup M_{\text{fill}}$ be the set of available moves (legal sprout sites and unfilled inner horns), and let $k : M \rightarrow \mathbb{R}_{>0}$ be a cost function depending only on a move's level d in the coherence tower.

Definition 3.22 (The inference problem). Find the probability distribution q on M maximizing the Shannon entropy $H(q) = -\sum_m q(m) \log q(m)$ subject to

- (C1) normalization, $\sum_m q(m) = 1$;
- (C2) budget, $\mathbb{E}_q[k] = R$;
- (C3) locality and homogeneity: $q(m)$ depends only on the isomorphism class of the local neighbourhood of m 's match.

Conjecture 3.23 (Gibbs rule). *The problem of Definition 3.22 has a unique solution of the form $q(m) \propto e^{-\lambda k(m)}$ with a single multiplier $\lambda(R)$. The sprout-to-fill ratio and the expansion law $N(t)$ are then determined by the level degeneracies (Catalan counts) and k , with no further choices.*

[GAP: Derivation is the subject of a dedicated session. Danger to watch: whether (C3) can be formulated without smuggling in a preferred scale.]

3.8 Coherence escalation and the growth law

Filling promotes questions up the associahedral tower (Section ??), where a question is an unfilled inner horn. A fill at level d mints questions at level $d + 1$, with supply growing like 4^d , so completion never depletes the supply of available moves and no region ever finishes.

The tower taxes fills alone. Sprouts have constant cost and remain available at every vertex, and by Remark 3.16 fills are metrically inert, so the growth of space is carried by sprouts and coherence escalation exerts no direct brake on it. The tower acts only through competition for budget, and under a Gibbs rule expensive moves lose share, so the mix shifts towards sprouting as the typical level rises.

The growth law is therefore decided by the second question of Remark 3.10. With one budget for the whole universe the move rate saturates at a constant once cheap sprouts dominate, vertex count grows linearly, and the scale factor decelerates. With a budget per region the total activity is proportional to the number of vertices, growth is exponential, and the late universe is a de Sitter phase. Locality favours the second reading. We assert neither, and we draw no comparison with the observed history of deceleration followed by late-time acceleration, since Remark 3.5 and problem O4 stand between vertex count and any scale factor. **[GAP: Settle the global versus local budget question and derive the growth law for at least one explicit cost function k ; compare against the toy simulation once the sampling estimator allows long runs.]**

Remark 3.24 (Two arrows). The framework contains two independent monotone quantities: the entropy S (Lemma 3.12) and the tower level d . That they are aligned is expected but unproved. **[GAP: prove or refute alignment]**

4 Results

5 Discussion

References

- [1] Jan Ambjørn, Bergfinnur Durhuus, and Jürg Fröhlich. Diseases of triangulated random surface models, and possible cures. *Nuclear Physics B*, 257:433–449, 1985.
- [2] Jan Ambjørn, Bergfinnur Durhuus, and Thordur Jonsson. *Quantum Geometry: A Statistical Field Theory Approach*. Cambridge University Press, 1997.
- [3] Jan Ambjørn, Andrzej Görlich, Jerzy Jurkiewicz, and Renate Loll. Nonperturbative quantum gravity. *Physics Reports*, 519:127–210, 2012.
- [4] Richard Arnowitt, Stanley Deser, and Charles W. Misner. The dynamics of general relativity. In Louis Witten, editor, *Gravitation: An Introduction to Current Research*, pages 227–265. Wiley, 1962.
- [5] Julian Barbour, Tim Koslowski, and Flavio Mercati. Identification of a gravitational arrow of time. *Physical Review Letters*, 113:181101, 2014.
- [6] Latham Boyle, Kieran Finn, and Neil Turok. CPT-symmetric universe. *Physical Review Letters*, 121:251301, 2018.
- [7] Graham Brightwell and Peter Winkler. Counting linear extensions. *Order*, 8:225–242, 1991.
- [8] François David. Planar diagrams, two-dimensional lattice gravity and surface models. *Nuclear Physics B*, 257:45–58, 1985.
- [9] Bryce S. DeWitt. Quantum theory of gravity. I. the canonical theory. *Physical Review*, 160:1113–1148, 1967.
- [10] Hartmut Ehrig, Karsten Ehrig, Ulrike Prange, and Gabriele Taentzer. *Fundamentals of Algebraic Graph Transformation*. Springer, 2006.
- [11] Marc Henneaux and Claudio Teitelboim. The cosmological constant and general covariance. *Physics Letters B*, 222:195–199, 1989.
- [12] Jacob Lurie. *Higher Topos Theory*. Princeton University Press, 2009.
- [13] Norman Margolus and Lev B. Levitin. The maximum speed of dynamical evolution. *Physica D*, 120:188–195, 1998.
- [14] Tullio Regge. General relativity without coordinates. *Il Nuovo Cimento*, 19:558–571, 1961.
- [15] Graeme Segal. Classifying spaces and spectral sequences. *Publications Mathématiques de l’IHÉS*, 34:105–112, 1968.
- [16] William G. Unruh. Unimodular theory of canonical quantum gravity. *Physical Review D*, 40:1048–1052, 1989.